

In [ ]:

```
%cd /content/drive/MyDrive/CGM
!ls
# !gzip "/content/drive/MyDrive/CGM/dataset/HIGGS_6M.csv.gz" -d "/content/drive/MyDrive/CGM/dataset"
```

In [ ]:

```
import xgboost
from numpy import loadtxt
from xgboost import XGBClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
```

In [ ]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

## EDA

In [ ]:


```
df=pd.read_csv("dataset/HIGGS_6M.csv")
df.head()
```


Out[ ]:

	1.0000000000000000e+00	8.692932128906250000e-01	6.35081827640534473e-01	2.256902605295181274e-01	3.2747006416320800e-01
0	1.0	0.907542	0.329147	0.359412	1.497
1	1.0	0.798835	1.470639	-1.635975	0.45
2	0.0	1.344385	-0.876626	0.935913	1.99
3	1.0	1.105009	0.321356	1.522401	0.88
4	0.0	1.595839	-0.607811	0.007075	1.81

In [ ]:

```
f = plt.figure(figsize=(10, 10))
plt.matshow(df.corr(), fignum=f.number)
cb = plt.colorbar()
cb.ax.tick_params(labelsize=14)
plt.title('Correlation Matrix', fontsize=16);
```









In [ ]:

```
dataset = pd.read_csv("dataset/HIGGS_6M.csv").to_numpy()
```


In [ ]:

```
for i in range(29):
    plt.hist(dataset[:,i],bins='auto')
    plt.title(str(i))
    plt.show()
```







3




4













5

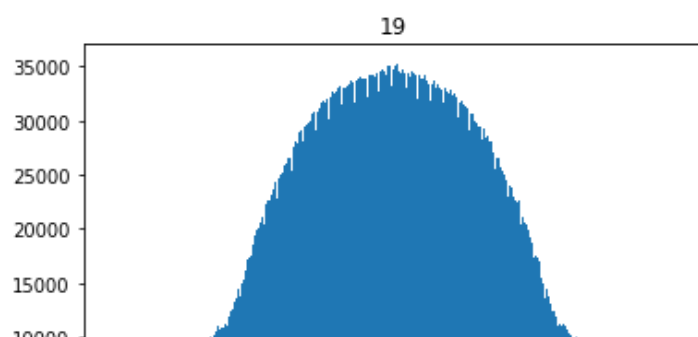











6
















```
-----
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-9-8e3c99999bc5> in <module>()
      1 for i in range(29):
----> 2     plt.hist(dataset[:,i],bins='auto')
      3     plt.title(str(i))
      4     plt.show()

/usr/local/lib/python3.7/dist-packages/matplotlib/pyplot.py in hist(x, bins, range, densi
ty, weights, cumulative, bottom, histtype, align, orientation, rwidth, log, color, label,
stacked, data, **kwargs)
    2608         align=align, orientation=orientation, rwidth=rwidth, log=log,
    2609         color=color, label=label, stacked=stacked, **({"data": data}
-> 2610         if data is not None else {}), **kwargs)
    2611
    2612

/usr/local/lib/python3.7/dist-packages/matplotlib/__init__.py in inner(ax, data, *args, *
**kwargs)
    1563     def inner(ax, *args, data=None, **kwargs):
    1564         if data is None:
-> 1565             return func(ax, *map(sanitize_sequence, args), **kwargs)
    1566
```



```

1567         bound = new_sig.bind(ax, *args, **kwargs)

/usr/local/lib/python3.7/dist-packages/matplotlib/axes/_axes.py in hist(self, x, bins, range, density, weights, cumulative, bottom, histtype, align, orientation, rwidth, log, color, label, stacked, **kwargs)
6727         patch = _barfunc(bins[:-1]+boffset, height, width,
6728                          align='center', log=log,
-> 6729                          color=c, **{bottom_kwarg: bottom})
6730         patches.append(patch)
6731         if stacked:

/usr/local/lib/python3.7/dist-packages/matplotlib/__init__.py in inner(ax, data, *args, **kwargs)
1563     def inner(ax, *args, data=None, **kwargs):
1564         if data is None:
-> 1565             return func(ax, *map(sanitize_sequence, args), **kwargs)
1566
1567         bound = new_sig.bind(ax, *args, **kwargs)

/usr/local/lib/python3.7/dist-packages/matplotlib/axes/_axes.py in bar(self, x, height, width, bottom, align, **kwargs)
2402         elif orientation == 'horizontal':
2403             r.sticky_edges.x.append(1)
-> 2404             self.add_patch(r)
2405             patches.append(r)
2406

/usr/local/lib/python3.7/dist-packages/matplotlib/axes/_base.py in add_patch(self, p)
1917         if p.get_clip_path() is None:
1918             p.set_clip_path(self.patch)
-> 1919             self._update_patch_limits(p)
1920             self.patches.append(p)
1921             p._remove_method = self.patches.remove

/usr/local/lib/python3.7/dist-packages/matplotlib/axes/_base.py in _update_patch_limits(self, patch)
1943             xys = patch_to_data.transform(xys)
1944
-> 1945             updatex, updatey = patch.get_transform().\
1946                 contains_branch_seperately(self.transData)
1947             self.update_datalim(xys, updatex=updatex,

/usr/local/lib/python3.7/dist-packages/matplotlib/patches.py in get_transform(self)
260     def get_transform(self):
261         """Return the `~.transforms.Transform` applied to the `Patch`."""
-> 262         return self.get_patch_transform() + artist.Artist.get_transform(self)
263
264     def get_data_transform(self):


/usr/local/lib/python3.7/dist-packages/matplotlib/patches.py in get_patch_transform(self)
775
776     def get_patch_transform(self):
-> 777         self._update_patch_transform()
778         return self._rect_transform
779

/usr/local/lib/python3.7/dist-packages/matplotlib/patches.py in _update_patch_transform(self)
754         """
755         x0, y0, x1, y1 = self._convert_units()
-> 756         bbox = transforms.Bbox.from_extents(x0, y0, x1, y1)
757         rot_trans = transforms.Affine2D()
758         rot_trans.rotate_deg_around(x0, y0, self.angle)

/usr/local/lib/python3.7/dist-packages/matplotlib/transforms.py in from_extents(*args)
787         The *y*-axis increases upwards.
788         """
-> 789         points = np.array(args, dtype=float).reshape(2, 2)
790         return Bbox(points)
791

```

KeyboardInterrupt:



In [ ]:

```
X = dataset[:,1:]
Y = dataset[:,0].astype(int)
print(X[0],Y[0])
print(np.shape(X), np.shape(X[0]), np.shape(Y), np.shape(Y[0]))
```

```
[ 9.07542109e-01  3.29147279e-01  3.59411865e-01  1.49796987e+00
 -3.13009530e-01  1.09553063e+00 -5.57524920e-01 -1.58822978e+00
  2.17307615e+00  8.12581182e-01 -2.13641927e-01  1.27101457e+00
  2.21487212e+00  4.99993950e-01 -1.26143181e+00  7.32156157e-01
  0.00000000e+00  3.98700893e-01 -1.13893008e+00 -8.19110195e-04
  0.00000000e+00  3.02219898e-01  8.33048165e-01  9.85699654e-01
  9.78098392e-01  7.79732168e-01  9.92355764e-01  7.98342586e-01] 1
(5999999, 28) (28,) (5999999,) ()
```

## Model on Raw dataset (no preprocessing)

You can skip training, jump to the last cell to load from drive

In [ ]:

```
seed = 7
test_size = 0.08
X_train, X_valid, y_train, y_valid = train_test_split(X, Y, test_size=test_size, random_
state=seed)
```

In [ ]:

```
eval_set = [(X_valid, y_valid)]
model = XGBClassifier()
model.fit(X_train, y_train, eval_metric="auc", eval_set=eval_set, verbose=True)
```

```
[20:38:34] WARNING: /workspace/src/learner.cc:686: Tree method is automatically selected
to be 'approx' for faster speed. To use old behavior (exact greedy algorithm on single ma
chine), set tree_method to 'exact'.
```

```
[0] validation_0-auc:0.679592
[1] validation_0-auc:0.684988
[2] validation_0-auc:0.688146
[3] validation_0-auc:0.692478
[4] validation_0-auc:0.708175
[5] validation_0-auc:0.71577
[6] validation_0-auc:0.723524
[7] validation_0-auc:0.728979
[8] validation_0-auc:0.730521
[9] validation_0-auc:0.731517
[10] validation_0-auc:0.736425
[11] validation_0-auc:0.736944
[12] validation_0-auc:0.740219
[13] validation_0-auc:0.743769
[14] validation_0-auc:0.745711
[15] validation_0-auc:0.745926
[16] validation_0-auc:0.747654
[17] validation_0-auc:0.748516
```

[17] validation\_0-auc:0.749516  
[18] validation\_0-auc:0.751581  
[19] validation\_0-auc:0.753539  
[20] validation\_0-auc:0.755207  
[21] validation\_0-auc:0.756212  
[22] validation\_0-auc:0.758195  
[23] validation\_0-auc:0.759694  
[24] validation\_0-auc:0.760278  
[25] validation\_0-auc:0.761302  
[26] validation\_0-auc:0.762495  
[27] validation\_0-auc:0.763558  
[28] validation\_0-auc:0.764722  
[29] validation\_0-auc:0.76556  
[30] validation\_0-auc:0.766682  
[31] validation\_0-auc:0.767155  
[32] validation\_0-auc:0.768227  
[33] validation\_0-auc:0.768552  
[34] validation\_0-auc:0.76924  
[35] validation\_0-auc:0.770104  
[36] validation\_0-auc:0.770717  
[37] validation\_0-auc:0.771381  
[38] validation\_0-auc:0.77207  
[39] validation\_0-auc:0.772801  
[40] validation\_0-auc:0.77311  
[41] validation\_0-auc:0.773857  
[42] validation\_0-auc:0.774087  
[43] validation\_0-auc:0.774715  
[44] validation\_0-auc:0.775203  
[45] validation\_0-auc:0.775696  
[46] validation\_0-auc:0.776102  
[47] validation\_0-auc:0.776552  
[48] validation\_0-auc:0.777135  
[49] validation\_0-auc:0.777567  
[50] validation\_0-auc:0.777888  
[51] validation\_0-auc:0.778765  
[52] validation\_0-auc:0.77896  
[53] validation\_0-auc:0.779252  
[54] validation\_0-auc:0.779619  
[55] validation\_0-auc:0.779822  
[56] validation\_0-auc:0.779892  
[57] validation\_0-auc:0.780381  
[58] validation\_0-auc:0.780899  
[59] validation\_0-auc:0.781338  
[60] validation\_0-auc:0.781614  
[61] validation\_0-auc:0.781701  
[62] validation\_0-auc:0.781829  
[63] validation\_0-auc:0.78201  
[64] validation\_0-auc:0.782281  
[65] validation\_0-auc:0.782884  
[66] validation\_0-auc:0.783448  
[67] validation\_0-auc:0.783688  
[68] validation\_0-auc:0.784032  
[69] validation\_0-auc:0.784123  
[70] validation\_0-auc:0.784336  
[71] validation\_0-auc:0.784749  
[72] validation\_0-auc:0.785088  
[73] validation\_0-auc:0.785265  
[74] validation\_0-auc:0.785427  
[75] validation\_0-auc:0.785509  
[76] validation\_0-auc:0.785724  
[77] validation\_0-auc:0.786082  
[78] validation\_0-auc:0.786251  
[79] validation\_0-auc:0.786409  
[80] validation\_0-auc:0.786623  
[81] validation\_0-auc:0.78675  
[82] validation\_0-auc:0.787107  
[83] validation\_0-auc:0.787251  
[84] validation\_0-auc:0.787678  
[85] validation\_0-auc:0.787871  
[86] validation\_0-auc:0.78826  
[87] validation\_0-auc:0.788387  
[88] validation\_0-auc:0.788568  
[89] validation\_0-auc:0.788700

```
[89] validation_0-auc:0.788789
[90] validation_0-auc:0.788916
[91] validation_0-auc:0.789111
[92] validation_0-auc:0.789268
[93] validation_0-auc:0.789334
[94] validation_0-auc:0.789485
[95] validation_0-auc:0.789628
[96] validation_0-auc:0.789814
[97] validation_0-auc:0.789979
[98] validation_0-auc:0.790176
[99] validation_0-auc:0.7904
```

Out[ ]:

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=1, gamma=0,
              learning_rate=0.1, max_delta_step=0, max_depth=3,
              min_child_weight=1, missing=None, n_estimators=100, n_jobs=1,
              nthread=None, objective='binary:logistic', random_state=0,
              reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
              silent=None, subsample=1, verbosity=1)
```

In [ ]:

```
y_pred = model.predict(X_valid)
predictions = [round(value) for value in y_pred]
accuracy = accuracy_score(y_valid, predictions)
print("Accuracy: %.2f%%" % (accuracy * 100.0))
```

Accuracy: 71.40%

In [ ]:

```
mod_dataset=dataset
logtransform=[1,4,6,10,14,18,22,23,24,25,26,27,28]
for index in logtransform:
    mod_dataset[:,index]=np.log(mod_dataset[:,index])
```

In [ ]:

```
for i in range(29):
    plt.hist(mod_dataset[:,i],bins='auto')
    plt.title(str(i))
    plt.show()
```

In [ ]:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
categorical=[0,9,13,17,21]
non_categorical=[]
for i in range(29):
    if i not in categorical:
        non_categorical.append(i)
print(non_categorical)
```

In [ ]:

```
for index in non_categorical:
    mod_dataset[:,index]=scaler.fit_transform(mod_dataset[:,index].reshape(-1,1)).reshape(-1)
```

In [ ]:

```
for i in range(29):
    plt.hist(mod_dataset[:,i],bins='auto')
    plt.title(str(i))
    plt.show()
```

In [ ]:

```
X = mod_dataset[:,1:1]
```

```
Y = mod_dataset[:,0].astype(int)
print(X[0],Y[0])
print(np.shape(X),np.shape(X[0]),np.shape(Y),np.shape(Y[0]))
```

```
[ 9.91932443e-02  3.26245081e-01  3.56807259e-01  8.94439464e-01
-3.11098036e-01  4.54615981e-01 -5.51916670e-01 -1.57885773e+00
 2.17307615e+00 -1.99794457e-01 -2.11991992e-01  1.26350201e+00
 2.21487212e+00 -1.24524704e+00 -1.25092734e+00  7.27702034e-01
 0.00000000e+00 -1.69832846e+00 -1.13024130e+00 -8.23977171e-04
 0.00000000e+00 -2.77900602e+00 -6.04251123e-01 -4.34355339e-01
 8.12042371e-02 -1.89601671e-01  2.57272610e-02 -5.25906146e-01] 1
(5999999, 28) (28,) (5999999,) ()
```

## Model on Processed Dataset

- To some features logarithm transform applied
- Standard Scaled
- Min Max Scaled
- Categorical features rounded to integers

You can skip training, jump to the last cell to load from drive

In [ ]:

```
seed = 7
test_size = 0.08
X_train, X_valid, y_train, y_valid = train_test_split(X, Y, test_size=test_size, random_
state=seed)
```

In [ ]:

```
eval_set = [(X_valid, y_valid)]
model = XGBClassifier()
eval_set = [(X_valid, y_valid)]
model.fit(X_train, y_train, eval_metric="auc", eval_set=eval_set, verbose=True)
```

In [ ]:

```
y_pred = model.predict(X_valid)
predictions = [round(value) for value in y_pred]
accuracy = accuracy_score(y_valid, predictions)
print("Accuracy: %.2f%%" % (accuracy * 100.0))
```

Accuracy: 71.38%


In [ ]:

```
from matplotlib import pyplot

results = model.evals_result()
print(results)
epochs = len(results['validation_0']['auc'])
x_axis = range(0, epochs)
# plot log loss
fig, ax = pyplot.subplots()
ax.plot(x_axis, results['validation_0']['auc'], label='Train')
ax.legend()
pyplot.ylabel('AUC')
pyplot.title('XGBoost AUC')
pyplot.show()
```

```
{'validation_0': {'auc': [0.679592, 0.68499, 0.687248, 0.691798, 0.708075, 0.709898, 0.72
2665, 0.723656, 0.729072, 0.733647, 0.734079, 0.737091, 0.740615, 0.74296, 0.745188, 0.74
7834, 0.748963, 0.750014, 0.752192, 0.753838, 0.754897, 0.756489, 0.758168, 0.759406, 0.7
60729, 0.762003, 0.763188, 0.763876, 0.765118, 0.765563, 0.76654, 0.767364, 0.768302, 0.7
68978, 0.769943, 0.770313, 0.771099, 0.771922, 0.772615, 0.773035, 0.773421, 0.773835, 0.
774293, 0.77491, 0.775118, 0.775384, 0.775844, 0.77643, 0.776753, 0.777219, 0.777556, 0.7
78001, 0.778301, 0.779054, 0.779275, 0.779472, 0.780207, 0.780469, 0.780745, 0.781102, 0.
78121, 0.781446, 0.781914, 0.78219, 0.782477, 0.782645, 0.782943, 0.783189, 0.783428, 0.7
```

```
83651, 0.783935, 0.784048, 0.784195, 0.784507, 0.784786, 0.784964, 0.785211, 0.785401, 0.785639, 0.785742, 0.785923, 0.786196, 0.78641, 0.78677, 0.78718, 0.787388, 0.78755, 0.787582, 0.787816, 0.788135, 0.788397, 0.788557, 0.788672, 0.789009, 0.789217, 0.78939, 0.789482, 0.78959, 0.789736, 0.789893]]}]
```



## Load Model from Drive

In [ ]:

```
import pickle
pickle.dump(model, open("xgb.pickle.dat", "wb"))
```

In [ ]:

```
"""
Checking whether load and save is working okay
"""
loaded_model = pickle.load(open("xgb.pickle.dat", "rb"))
y_pred = loaded_model.predict(X_valid)
predictions = [round(value) for value in y_pred]
accuracy = accuracy_score(y_valid, predictions)
print("Accuracy: %.2f%%" % (accuracy * 100.0))
```

In [ ]:

```
"""
too slow abandoned after 10 minutes
"""
# df = pd.DataFrame(dataset, columns=np.arange(29)).sample(10)
# sns.pairplot(df)
```

In [ ]:

```
"""
took too long to run, abandoned after 14 minutes
"""
# from sklearn.neighbors import KNeighborsClassifier
# from sklearn import metrics
# knn=KNeighborsClassifier(n_neighbors=2)
# knn.fit(X_train,y_train)
# y_pred=knn.predict(X_valid)
# score=metrics.accuracy_score(y_valid,y_pred)
# print(score)
```